

Final Project_Mouseketeers

Three Mouseketeers (Co)

2025-12-10

```
library(caret)

## Loading required package: ggplot2

## Loading required package: lattice

library(glmnet)

## Loading required package: Matrix

## Loaded glmnet 4.1-10

install.packages("pROC", repos = "https://cran.rstudio.com/")

##
## The downloaded binary packages are in
## /var/folders/m0/q7ycmrxr14311yyb1w1lfh5680000gn/T//RtmpxdaYqV downloaded_packages

library(pROC)

## Type 'citation("pROC")' for a citation.

##
## Attaching package: 'pROC'

## The following objects are masked from 'package:stats':
## 
##     cov, smooth, var

library(ggplot2)
```

1. Compiled dataframe to include both summaries of glucose reading and patient demographic data.

#Downloading the dataframe by Connor, compiled from multiple data frames, including variables
MasterFrame <- read.csv("~/Desktop/Final Project/combined_patient_data.csv")

2. Based on literature review and domain knowledge, identify percent out of range as the predicted value (ideal range is between 70 and 180).

```

#Connor was able to identify percent out of range in the pct_70_180 column.
#I will create a new column "Uncontrolled" which is a binary variable.
#It will be set as 1 if it is uncontrolled and 0 if it is controlled.
#Uncontrolled is defined as pct_70_180 >70 (%)
MasterFrame$Uncontrolled <- ifelse(MasterFrame$pct_70_180 < 70, 1, 0)

```

3. Based on literature review and domain knowledge, include these variables as predictors

Connor has chosen these variables based on his expertise as an endocrinologist. (predictor_vars)

```

predictor_vars <- c(
  # Demographics
  "AgeAtEnrollment",
  "Gender",
  "Ethnicity",
  "Race",
  "DiagAge",
  "DiagAgeApprox",
  "Weight_kg",
  "Height_cm",
  "BMI",
  "BldPrSys",
  "BldPrDia",
  "EducationLevel",
  "AnnualIncome",
  "InsuranceType",

  # Baseline clinical data
  "HbA1c_Screening",
  "NumMedicalConditions",
  "NumMedications",

  # TDD metrics
  "tdd_mean",
  "tdd_median",
  "tdd_sd",
  "tdd_cv",
  "TDD_per_kg",
  "TDD_per_BMI",

  # Basal insulin metrics
  "basal_mean",
  "basal_median",
  "basal_sd",
  "basal_pct",

  # Bolus insulin metrics
  "bolus_mean",
  "bolus_median",
  "bolus_sd",
  "bolus_pct",
  "mean_boluses_per_day",

```

```

# Data collection duration
"n_days"
)

```

I have some concerns regarding the use of tdd and insulin metrics, as it may be a source of leakage for prediction models. (insulin_vars)

```

insulin_vars <- c(
  "tdd_mean", "tdd_median", "tdd_sd", "tdd_cv", "TDD_per_kg", "TDD_per_BMI",
  "basal_mean", "basal_median", "basal_sd", "basal_pct",
  "bolus_mean", "bolus_median", "bolus_sd", "bolus_pct",
  "mean_boluses_per_day"
)

```

This chunk of code will narrow down the table to only include these variables and the Uncontrolled outcome variable.

```

ModelDF_full <- MasterFrame[, c(predictor_vars, "Uncontrolled")]
ModelDF_no_insulin <- MasterFrame[, c(setdiff(predictor_vars, insulin_vars), "Uncontrolled")]

#Drop observations with no outcome
ModelDF_full_outcome <- ModelDF_full[!is.na(ModelDF_full$Uncontrolled), ]
ModelDF_no_insulin_outcome <- ModelDF_no_insulin[!is.na(ModelDF_no_insulin$Uncontrolled), ]

clean_impute <- function(df) {
  #Impute missing predictor variables as median
  preProc <- preProcess(df, method = "medianImpute")
  df <- predict(preProc, df)
  #Change characters as factors (to prepare for glm)
  char_cols <- sapply(df, is.character)
  df[, char_cols] <- lapply(df[, char_cols], as.factor)

  return(df)
}

ModelDF_full_clean <- clean_impute(ModelDF_full_outcome)
ModelDF_no_insulin_clean <- clean_impute(ModelDF_no_insulin_outcome)

```

4. Split dataframes to Train/Test

```

set.seed(1)

# library(caret) already downloaded above
# Partition function -> To ensure acceptable ratio of controlled/uncontrolled in train/test sets

#Split Full set into train and test with 80/20 ratio
train_index_full <- createDataPartition(ModelDF_full_clean$Uncontrolled, p = 0.8, list = FALSE)
Train_full <- ModelDF_full_clean[train_index_full, ]
Test_full <- ModelDF_full_clean[-train_index_full, ]

# Same for No Insulin Set
train_index_no_insulin <- createDataPartition(ModelDF_no_insulin_clean$Uncontrolled, p = 0.8, list = FALSE)

```

```

Train_no_insulin <- ModelDF_no_insulin_clean[train_index_no_insulin, ]
Test_no_insulin <- ModelDF_no_insulin_clean[-train_index_no_insulin, ]

# Optional: check proportions
prop.table(table(Train_full$Uncontrolled))

## 
##          0           1
## 0.56666667 0.4333333

prop.table(table(Test_full$Uncontrolled))

## 
##          0           1
## 0.6363636 0.3636364

prop.table(table(Train_no_insulin$Uncontrolled))

## 
##          0           1
## 0.5888889 0.4111111

prop.table(table(Test_no_insulin$Uncontrolled))

## 
##          0           1
## 0.5454545 0.4545455

```

5. Ridge/Lasso to narrow down variables to 6 most relevant to prediction

Options for variable selection include: Subset selection (not good for variables that might have multicollinearity), ridge/lasso (shrinking some variables close to zero and zero respectively) and principle component analysis (PCA, which creates new variables which are amalgamations of other variables, not very interpretable).

Ridge/Lasso is a good balance between interpretability and handling collinearity.

```

# library(glmnet) downloaded above

# setting the predictors (x) and outcome (y)
x <- model.matrix(Uncontrolled ~ ., data = Train_full)[,-1]
y <- Train_full$Uncontrolled

# Cross-validation to find best lambda
cv.lasso <- cv.glmnet(x, y, alpha = 1, family = "binomial")
best_lambda_lasso <- cv.lasso$lambda.min
best_lambda_lasso

## [1] 0.06909099

```


```

# Cross-validation for best lambda
cv.ridge <- cv.glmnet(x, y, alpha = 0, family = "binomial")
best_lambda_ridge <- cv.ridge$lambda.min
best_lambda_ridge

## [1] 0.2159588

# Lasso coefficient path
lasso.mod <- glmnet(x, y, alpha = 1, family = "binomial")
plot(lasso.mod, xvar = "lambda", label = TRUE, main = "Lasso Coefficient Paths")


```



```

# Ridge coefficient path
ridge.mod <- glmnet(x, y, alpha = 0, family = "binomial")
plot(ridge.mod, xvar = "lambda", label = TRUE, main = "Ridge Coefficient Paths")

```



```
#I ultimately chose to use lasso, which sets other variables to zero
coef_lasso <- predict(lasso.mod, type = "coefficients", s = best_lambda_lasso)
nonzero_vars_full <- rownames(coef_lasso)[coef_lasso[,1] != 0]
print(nonzero_vars_full)
```

```
## [1] "(Intercept)"           "HbA1c_Screening"      "basal_sd"
## [4] "mean_boluses_per_day"
```


Do the same for the no insulin set to see if there are the same predictors

```
x_no_insulin <- model.matrix(Uncontrolled ~ ., data = Train_no_insulin)[,-1]
y_no_insulin <- Train_no_insulin$Uncontrolled


# Lasso
cv.lasso_no_insulin <- cv.glmnet(x_no_insulin, y_no_insulin, alpha = 1, family = "binomial")
best_lambda_lasso_no_insulin <- cv.lasso_no_insulin$lambda.min

# Ridge
cv.ridge_no_insulin <- cv.glmnet(x_no_insulin, y_no_insulin, alpha = 0, family = "binomial")
best_lambda_ridge_no_insulin <- cv.ridge_no_insulin$lambda.min

# Visualization
lasso.mod_no_insulin <- glmnet(x_no_insulin, y_no_insulin, alpha = 1, family = "binomial")
plot(lasso.mod_no_insulin, xvar = "lambda", label = TRUE, main = "Lasso Coefficient Paths (No Insulin)")
```



```
ridge.mod_no_insulin <- glmnet(x_no_insulin, y_no_insulin, alpha = 0, family = "binomial")
plot(ridge.mod_no_insulin, xvar = "lambda", label = TRUE, main = "Ridge Coefficient Paths (No Insulin)")
```



```
# Get nonzeros from lasso with best lambda
coef_lasso_no_insulin <- predict(lasso.mod_no_insulin, type = "coefficients", s = best_lambda_lasso_no_
nonzero_vars_no_insulin <- rownames(coef_lasso_no_insulin)[coef_lasso_no_insulin[,1] != 0]
print(nonzero_vars_no_insulin)
```

```
## [1] "(Intercept)"      "AgeAtEnrollment" "HbA1c_Screening"
```

6. Create a glm model based on the selected variables using the train data set

```
glm_model_full <- glm(Uncontrolled ~
                        HbA1c_Screening +
                        basal_sd +
                        mean_boluses_per_day,
                        data = Train_full,
                        family = binomial)

summary(glm_model_full)

##
## Call:
## glm(formula = Uncontrolled ~ HbA1c_Screening + basal_sd + mean_boluses_per_day,
##      family = binomial, data = Train_full)
##
## Coefficients:
```

```

##                               Estimate Std. Error z value Pr(>|z|)
## (Intercept)           -22.5882     5.4244 -4.164 3.13e-05 ***
## HbA1c_Screening       1.9522     0.5453  3.580 0.000344 ***
## basal_sd              0.6073     0.1539  3.945 7.98e-05 ***
## mean_boluses_per_day  0.4621     0.1957  2.361 0.018235 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
## Null deviance: 123.162  on 89  degrees of freedom
## Residual deviance: 63.387  on 86  degrees of freedom
## AIC: 71.387
##
## Number of Fisher Scoring iterations: 6

```

Same for no insulin

```

glm_model_no_insulin <- glm(Uncontrolled ~ AgeAtEnrollment +
                                HbA1c_Screening,
                                data = Train_no_insulin,
                                family = binomial)

summary(glm_model_no_insulin)

```

```

##
## Call:
## glm(formula = Uncontrolled ~ AgeAtEnrollment + HbA1c_Screening,
##      family = binomial, data = Train_no_insulin)
##
## Coefficients:
##                               Estimate Std. Error z value Pr(>|z|)
## (Intercept)           -10.22352   2.87993 -3.550 0.000385 ***
## AgeAtEnrollment      -0.02915   0.01747 -1.669 0.095134 .
## HbA1c_Screening        1.41895   0.36472  3.891 0.000100 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
## Null deviance: 121.907  on 89  degrees of freedom
## Residual deviance: 86.486  on 87  degrees of freedom
## AIC: 92.486
##
## Number of Fisher Scoring iterations: 5

```

7. Cross validate Full

```

# 1. Predict probabilities on Test_full
prob_full <- predict(glm_model_full, newdata = Test_full, type = "response")

# 2. Convert probabilities to binary predictions (threshold 0.5)
pred_full <- ifelse(prob_full > 0.5, 1, 0)

```

```

# 3. Create Confusion Matrix
conf_matrix_full <- table(Actual = Test_full$Uncontrolled, Predicted = pred_full)
print(conf_matrix_full)

##          Predicted
## Actual   0   1
##       0 10  4
##       1  2  6

accuracy_full <- sum(diag(conf_matrix_full)) / sum(conf_matrix_full) * 100
print(accuracy_full)

## [1] 72.72727

# 1. Predict probabilities on Test_no_insulin
prob_no_insulin <- predict(glm_model_no_insulin, newdata = Test_no_insulin, type = "response")

# 2. Convert probabilities to binary predictions (threshold 0.5)
pred_no_insulin <- ifelse(prob_no_insulin > 0.5, 1, 0)

# 3. Create Confusion Matrix
conf_matrix_no_insulin <- table(Actual = Test_no_insulin$Uncontrolled, Predicted = pred_no_insulin)
print(conf_matrix_no_insulin)

##          Predicted
## Actual   0   1
##       0  9  3
##       1  5  5

accuracy_no_insulin <- sum(diag(conf_matrix_no_insulin)) / sum(conf_matrix_no_insulin) * 100
print(accuracy_no_insulin)

## [1] 63.63636

```

8. Creating AUC


```

# already installed library(pROC)
roc_full <- roc(response = Test_full$Uncontrolled, predictor = prob_full, quiet = TRUE)
plot(roc_full,
     col = "blue",
     lwd = 3,
     main = "ROC Curve: Full Model (Test Set)",
     xlab = "False Positive Rate (1 - Specificity)",
     ylab = "True Positive Rate (Sensitivity)")

# Add the diagonal line (random guess)
abline(a = 1, b = -1, lty = 2, col = "gray50")

text(x = 0.5, y = 0.4,
      labels = paste("AUC =", round(auc(roc_full), 3)),
      col = "blue",
      cex = 1.2)

```



```

roc_no_insulin <- roc(response = Test_no_insulin$Uncontrolled, predictor = prob_no_insulin, quiet = TRUE)

plot(roc_no_insulin,
      col = "red",
      lwd = 3,
      main = "ROC Curve: No-Insulin Model (Test Set)",
      xlab = "False Positive Rate (1 - Specificity)",
      ylab = "True Positive Rate (Sensitivity)")

# Add the diagonal line (random guess)
abline(a = 1, b = -1, lty = 2, col = "gray50")

# Add AUC as text on the plot
text(x = 0.5, y = 0.4,
      labels = paste("AUC =", round(auc(roc_no_insulin), 3)),
      col = "red",
      cex = 1.2)

```

ROC Curve: No-Insulin Model (Test Set)

